

Module 0: Welcome to STA 310

Rebecca C. Steorts (slide and course adaptation from Maria Tackett)

Welcome!

Teaching Team

Instructor:

Professor Rebecca Steorts
Old Chem 208
rebecca.steorts@duke.edu

Teaching assistants

Devarpita (Deva) Bag,
PhD Student
devarpita.bag@duke.edu

Announcements

- ▶ We will meet on Friday, January 9th for lecture, so please come prepared for lecture and not lab (as the PhD student is traveling and not available for lab). I will be sure to give you this lecture back at the end of the semester.
- ▶ The course webpage (<https://resteorts.github.io/teach/generalized.html>) will be updated on roughly a weekly basis, so please check this frequently for any updates.

Course logistics

Lectures

Tuesday and Thursday, 11:45 - 1:00 pm, Perkins 127

Labs (Office Hour or Alternate Lecture Time)

Lab 01: Friday, 11:45 - 1:00 pm, Old Chemistry 001

Generalized Linear Models

In statistics, a generalized linear model (GLM) is a flexible generalization of ordinary linear regression. The GLM generalizes linear regression by allowing the linear model to be related to the response variable via a link function and by allowing the magnitude of the variance of each measurement to be a function of its predicted value.¹

¹Source: Generalized linear model

Generalized Linear Models

In statistics, a generalized linear model (GLM) is a flexible generalization of ordinary linear regression. The GLM generalizes linear regression by allowing the linear model to be related to the response variable via a link function and by allowing the magnitude of the variance of each measurement to be a function of its predicted value.¹

Example: Logistic regression

$$\begin{aligned}\pi = P(y = 1|x) &\Rightarrow \text{Link function: } \log\left(\frac{\pi}{1 - \pi}\right) \\ &\Rightarrow \log\left(\frac{\pi}{1 - \pi}\right) = \beta_0 + \beta_1 \times\end{aligned}$$

¹Source: Generalized linear model

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.
- ▶ explain how specific models fit into the GLM framework

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.
- ▶ explain how specific models fit into the GLM framework
- ▶ identify the appropriate model given the data and analysis objective.

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.
- ▶ explain how specific models fit into the GLM framework
- ▶ identify the appropriate model given the data and analysis objective.
- ▶ analyze real-world data by fitting and interpreting GLMs.

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.
- ▶ explain how specific models fit into the GLM framework
- ▶ identify the appropriate model given the data and analysis objective.
- ▶ analyze real-world data by fitting and interpreting GLMs.
- ▶ use R for analysis and write reports

Course learning objectives

By the end of the semester, you will be able to ...

- ▶ describe generalized linear models (GLMs) as a unified framework.
- ▶ explain how specific models fit into the GLM framework
- ▶ identify the appropriate model given the data and analysis objective.
- ▶ analyze real-world data by fitting and interpreting GLMs.
- ▶ use R for analysis and write reports
- ▶ effectively communicate results from statistical analyses to a general audience in writing.

Course topics

Generalized Linear Models

- ▶ Review of distributions, likelihoods, and regression
- ▶ Introduce models for non-normal response variables
- ▶ Estimation, interpretation, and inference
- ▶ Mathematical details of GLMs as a unified framework

GLMs in practice

Papers

Officiating bias: The effect of foul differential on foul calls in NCAA basketball

Kyle J. Anderson & David A. Pierce

Pages 687-694 | Accepted 07 Jan 2009, Published online: 20 May 2009

Cite this article <https://doi.org/10.1080/02640410902729733>

Full Article

Figures & data

References

Citations

Metrics

Reprints & Permissions

Read this article

" . . . a **logistic regression model** is used to test how the likelihood of a foul is affected by which team is the home team, the foul differential, and the score differential. . . The logistic regression was run under several specifications . . . using **clustered observation standard errors**, with each game as a cluster. This is done as an attempt to adjust for the fact that **observations may not be independent** as required under the logistic specification.²

²Anderson, K. J., & Pierce, D. A. (2009). Officiating bias: The effect of foul differential on foul calls in NCAA basketball. *Journal of sports sciences*, 27(7), 687-694.

GLMs in practice



JSLHR ■

Research Article

Intersectionality of Race and Question-Asking in Women After Right Hemisphere Brain Damage

Danai Kasambira Fannin,^a Jada Elleby,^a Maria Tackett,^b and Jamila Minga^c

^aDepartment of Communication Sciences and Disorders, North Carolina Central University, Durham ^bDepartment of Statistical Science, Duke University, Durham, NC ^cDepartment of Head and Neck Surgery & Communication Sciences and Department of Neurology, Vascular and Stroke Division, Duke University School of Medicine, Durham, NC

“... we used **negative binomial regression** to model the association between the number of questions produced, race, and group after adjusting for the additional covariates age and years of education. **Poisson and zero-inflated Poisson regression models** were also considered... the negative binomial model was a good fit for the data given the **overdispersion** in the distribution of number of questions asked.”³

³Fannin, D. K., Elleby, J., Tackett, M., & Minga, J. (2023). Intersectionality of Race and Question-Asking in Women After Right Hemisphere Brain Damage. *Journal of Speech, Language, and Hearing Research, 66*(1), 314-324.

GLMs in practice (continued)

- ▶ The authors consider data that follows a Poisson distribution.
- ▶ Recall that the Poisson distribution has the same mean and variance parameter (λ).
- ▶ Due to this, it is possible for overdispersion to occur meaning that then there is more variation in the response than what's implied by a Poisson model.
- ▶ This means the standard errors of the model coefficients are artificially small.

This implies the following:

- ▶ The p-values are artificially small
- ▶ Can lead to models that are more complex than what is needed

As such, we will need to consider a careful treatment of this (and this is what the authors did in the paper above).

Meet your classmates!

- ▶ Get in groups of 2 - 3
- ▶ Each person in the group...
 - ▶ Introduce yourself (name, major, year of study, hometown)
 - ▶ Share something interesting about yourself
- ▶ Everyone will introduce one person from your group to the class

Course details

Pre-reqs

Pre-reqs

STA 210 and STA 230 / STA 240

Background knowledge

Statistical methods

- ▶ Linear and logistic regression
- ▶ Statistical inference
- ▶ Basic understanding of random variables

Computing

- ▶ Using R for data analysis
- ▶ Writing reports using Rmd
- ▶ Understanding of github
- ▶ Understanding reproducibility

Course toolkit

Course webpage:

<https://resteorts.github.io/teach/generalized.html>

- ▶ Course information and course schedule

Canvas

- ▶ Changes to Schedule
- ▶ Ed Discussion
- ▶ Homework uploads

Gradescope (link on course webpage)

- ▶ Homework uploads (make sure to upload to Canvas as well).

Ed Discussion (link on course webpage)

- ▶ Course discussion

Class Meetings

Lectures

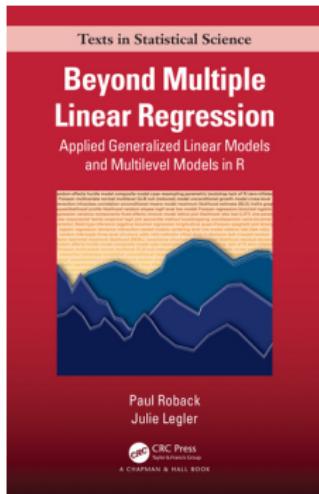
- ▶ Some traditional lecture
- ▶ Short individual and group activities
- ▶ Bring fully-charged laptop / tablet to use R

Labs (start January 9)

- ▶ Work on class assignments with TA support
- ▶ Time for clarifying questions regarding course material
- ▶ **Alternative lecture time** when needed (such as this week)

Attendance is strongly expected (if you are healthy!)

Readings



- ▶ Primary textbook: *Beyond Multiple Linear Regression* by Roback and Legler
- ▶ Other texts:
 - ▶ *R for Data Science (2nd edition)* by Wickham, Çetinkaya-Rundel, and Grolemund
 - ▶ Tidy Modeling with R by Kuhn and Silge
 - ▶ Articles and videos periodically assigned

Computing toolkit

R and RStudio

- ▶ Install R and RStudio on your laptop
- ▶ Click [here](#) for instructions to install RStudio and configure git

or

Access RStudio through Docker container provided by Duke OIT

- ▶ Reserve a generic **RStudio** container (there is no course specific container)

Canvas and Gradescope

- ▶ All homework assignments will be uploaded to Gradescope and Canvas.
- ▶ Gradescope allows more fair and balanced grading.
- ▶ Canvas allows us to check the reproducibility of your work.
- ▶ Unfortunately, there is no platform that does both (to my knowledge).
- ▶ Feedback will be given in Gradescope and is individual and private.

Ed Discussion

- ▶ Online discussion forum (like Piazza, etc.)
- ▶ Platform to ask questions about course content, logistics, assignments, etc.
- ▶ Content organized by channels. Before posting, please browse previous posts to see if your question has already been answered. If not, please post your question in the relevant channel.
- ▶ Questions about grades, absences, and other private matters should be emailed to me with “STA 310” in the subject line.

Activities & Assessment

Homework (40%)

- ▶ Individual assignments
- ▶ Combination of conceptual questions, guided analyses, and open-ended analyses

Quizzes (60%)

- ▶ Individual online quizzes
- ▶ Covers content since the previous quiz, including readings, lecture notes, in-class activities, and homework

Grading

Final grades will be calculated as follows

Category	Percentage
Homework	40%
Quizzes	60%

See the course syllabus for letter grade thresholds.

Course community

Course community

- ▶ Uphold the Duke Community Standard:
 - ▶ *I will not lie, cheat, or steal in my academic endeavors;*
 - ▶ *I will conduct myself honorably in all my endeavors;*
 - ▶ *I will act if the Standard is compromised.*
- ▶ Commit to respect, honor, and celebrate our diverse community
- ▶ Commit to being part of a learning environment that is welcoming and accessible to everyone

Accessibility

- ▶ The Student Disability Access Office (SDAO) is available to ensure that students are able to engage with their courses and related assignments.
- ▶ If you have documented accommodations from SDAO, please send the documentation within the first week to make sure all accommodations can be put in place as quickly as possible!
- ▶ I am committed to making all course activities and materials accessible. If any course component is not accessible to you in any way, please don't hesitate to let me know.

Support

- ▶ **Office hours** to meet with a member of the teaching team.
 - ▶ Find the course schedule on the course webpage
 - ▶ Office hours begin January 16
 - ▶ Please see me after class if you have questions before then.
- ▶ **Ed Discussion** for questions about course logistics, content, and assignments
- ▶ **Email** for questions not appropriate for Ed Discussion, e.g., regarding personal matters or grades
 - ▶ Please put **STA 310** in the subject line

See the syllabus regarding additional academic and mental health and wellness resources.

Latex Resources

1. <https://wch.github.io/latexsheet/latexsheet.pdf>
2. <https://www.bu.edu/math/files/2013/08/LongTeX1.pdf>
3. <https://www.docx2latex.com/tutorials/mathematical-equations-latex/>

Course Content

The course will closely follow the course textbook. Everyone is expected to follow the text readings.

1. Review of Reproducibility (Markdown versus Quarto)
2. Review of Multiple Linear Regression
3. Commonly Used Distributions (Bernoulli, Poisson, Gamma, Beta, others.)
4. Properties of Likelihoods
5. Poisson regression
6. The theory of exponential families and generalized linear models
7. Logistic regression
8. Multinomial regression
9. Advanced topics (as time permits)

The goal is to provide you with foundational applied and theoretical knowledge on GLMs.

Questions

Questions?